

✓ Training Strategy Refactor – To-Do List for Copilot Agent

Goal: **Keep the current separate training for Scenario 1 & 2, but:**

- Fix validation / evaluation so they work correctly with separate or unified logic.
- Add a **config toggle** to switch between:
 - `train_mode = "separate"` → two pipelines (S1 & S2)
 - `train_mode = "unified"` → single global pipeline

This file is designed as a step-by-step task list for an AI agent (e.g. Copilot).

1. Add a `train_mode` Toggle in Config

- Open the main configuration file (e.g. `config.yaml` or `config.py`).
- Add a new top-level key:
 - `train_mode: "separate"` # options: "separate", "unified"
- If using nested config sections, ensure both training scripts can access:
 - `config.train_mode`
- Document the allowed values and behavior in comments or README.

Example YAML snippet:

```
training:
  train_mode: "separate"  # "separate" or "unified"
```

2. Centralize Scenario Definitions

- Create a central module (e.g. `src/config/scenarios.py`) to define:
 - Scenario 1:
 - `horizon: months 0–23`
 - `available features: pre-LOE only`
 - Scenario 2:
 - `horizon: months 6–23`
 - `available features: pre-LOE + months 0–5 info`
- Expose helper functions:
 - `get_scenario_definition(name: str) -> dict`
 - `is_month_in_scenario(month, scenario_name)`
- Update training code to import scenario definitions from this module instead of hard-coding ranges in multiple places.

3. Keep Current Separate Training Pipelines (S1 & S2)

- Identify the **current** S1 training script/function (e.g. `train_scenario1()`):
 - Confirm it:
 - Uses pre-LOE features only.
 - Trains on months 0–23.
 - Produces S1 predictions for months 0–23.
 - Identify the **current** S2 training script/function (e.g. `train_scenario2()`):
 - Confirm it:
 - Uses pre-LOE + months 0–5 features.
 - Trains on months 6–23.
 - Produces S2 predictions for months 6–23.
 - Wrap them in a higher-level function `train_separate(config)` :
 - Calls `train_scenario1(config)` .
 - Calls `train_scenario2(config)` .
 - Stores models, metrics, and artifacts in separate subfolders:
 - `models/scenario1/`
 - `models/scenario2/`
 - `metrics/scenario1/`
 - `metrics/scenario2/`
-

4. Implement a Unified Training Pipeline (Single Model)

- Add a new function `train_unified(config)` :
 - Builds a **single** multi-horizon model:
 - Training rows: all months 0–23 for all brands.
 - Features include `months_postgx` and time-window flags (early / mid / late).
 - Uses a unified feature pipeline that can support both:
 - Scenario 1 (no early post-LOE info at origin).
 - Scenario 2 (has early post-LOE summary at origin).
 - Add a `scenario_flag` feature:
 - `scenario_flag = 0` for Scenario-1-like rows.
 - `scenario_flag = 1` for Scenario-2-like rows.
 - In training, create two types of samples:
 - S1-style: origin at month 0, features limited to pre-LOE info.
 - S2-style: origin at month 6, features include summary of months 0–5.
 - Concatenate S1-style and S2-style training rows into a single dataset.
 - Train **one global model** that can handle both (`scenario_flag` differentiates regimes).
-

5. Connect `train_mode` to the Main Entry Point

- In the main training entry script (e.g. `train.py`):

- Load config.
- Read `config.training.train_mode` (or equivalent).
- Use conditional logic:

```
if config.training.train_mode == "separate":  
    train_separate(config)  
elif config.training.train_mode == "unified":  
    train_unified(config)  
else:  
    raise ValueError(f"Unknown train_mode:  
{config.training.train_mode}")
```

- Ensure that all CLI entry points or notebooks honor this setting.
-

6. Fix / Unify Validation & Metric Code

6.1 Centralize Metric Computation

- Create a module `src/metrics/loe_metrics.py` with functions:
 - `compute_metric_scenario1(y_true, y_pred, avg_j, buckets)`
 - `compute_metric_scenario2(y_true, y_pred, avg_j, buckets)`
- Move any scattered metric code into these functions.
- Ensure the functions:
 - Use the official formulas for PE (S1 and S2).
 - Apply Bucket 1 and Bucket 2 weights correctly.
- Add unit tests for both metrics using small synthetic examples.

6.2 Validation for Separate Training

- In **separate mode** (`train_mode = "separate"`):
 - For Scenario 1:
 - Perform CV using only S1 pipeline and S1 metric.
 - For Scenario 2:
 - Perform CV using only S2 pipeline and S2 metric.
- Ensure validation outputs are stored as:
 - `metrics/val_scenario1.json`
 - `metrics/val_scenario2.json`

6.3 Validation for Unified Training

- In **unified mode** (`train_mode = "unified"`):
 - Use **GroupKFold by brand** for CV.
 - For each fold:
 - Generate predictions for all brands/months 0–23.

- Compute:
 - S1 metric on that fold.
 - S2 metric on that fold.
 - Aggregate metrics across folds:
 - `mean_PE_s1`
 - `mean_PE_s2`
 - A combined score, e.g. `0.5 * PE_s1 + 0.5 * PE_s2`.
 - Save unified validation outputs as:
 - `metrics/val_unified.json` (including both S1 and S2 metrics).
-

7. Ensure Feature Pipelines Behave Under Both Modes

- Identify the shared **feature engineering** module (e.g. `src/features/build_features.py`).
 - Add support for:
 - `scenario_flag` feature.
 - Optional early post-LOE summary features (mean 0–5, slope 0–5, last value at 5, etc.).
 - In **separate mode**:
 - For Scenario 1:
 - Disable early-post features (fill with 0 or NaN).
 - For Scenario 2:
 - Enable early-post features.
 - In **unified mode**:
 - Always build both S1-style and S2-style rows.
 - Ensure:
 - S1-style rows have `scenario_flag = 0`, early-post features disabled.
 - S2-style rows have `scenario_flag = 1`, early-post features computed.
-

8. Model Saving & Loading Conventions

- For **separate mode**:
 - Save models as:
 - `models/scenario1/model.pkl`
 - `models/scenario2/model.pkl`
 - For **unified mode**:
 - Save model as:
 - `models/unified/model.pkl`
 - Adjust prediction scripts to:
 - Load the correct model depending on `train_mode`.
 - Produce predictions sliced appropriately:
 - S1: months 0–23.
 - S2: months 6–23.
-

9. CLI / Notebook Documentation

- Update `README.md` or `docs/training.md` to explain:

- The purpose of `train_mode`.
 - Differences between `separate` and `unified` modes.
 - How to run each mode, e.g.:
 - `python train.py --train_mode separate`
 - `python train.py --train_mode unified`
 - Provide example commands for:
 - Training.
 - Validation only.
 - Generating S1 and S2 submissions from each mode.
-

10. Sanity Checks

- In both `separate` and `unified` modes:
 - Confirm that S1 predictions always have months 0–23 per brand.
 - Confirm that S2 predictions always have months 6–23 per brand.
 - Run a small smoke test with 2–3 brands to ensure the pipelines don't mix up horizons.
- Verify that validation metrics are consistent between:
 - Legacy (old) separate pipelines.
 - New refactored separate pipelines.
- Confirm that unified mode produces comparable or better metrics when hyperparameters are tuned.