

# novartis\_datathon\_2025-Arman – Improvement TODO for Copilot Agent

---

You are Copilot in VS Code, working inside the `novartis_datathon_2025-Arman` repo.

Goal: Turn the project from “complex and good” into “sharp, reliable, and leaderboard-dangerous.”

---

## 1. Config, Metrics & Validation Plumbing

### 1.1 Make Config the Single Source of Truth

- Find all places where scenario, model, sample weights, or feature choices are **hard-coded** in Python (especially in `train.py`, `inference.py`, `models/*.py`).
- Move those settings into YAML (e.g. `configs/run_defaults.yaml`, `configs/model_*.yaml`).
- Ensure all entry points load config via a shared function (e.g. `load_run_config()`), and do **not** silently override YAML settings in code.
- Add comments in YAML explaining each important flag (scenario, sample weights, model type).

### 1.2 Centralize Official Metric Usage

- Create or clean up a single module, e.g. `src/metrics/official.py`, that:
  - Wraps `compute_metric1` (Scenario 1) and `compute_metric2` (Scenario 2).
  - Provides helper functions:
    - `evaluate_on_train(df_train, df_pred_s1, df_pred_s2, df_aux)`
    - `evaluate_cv_fold(...)` (if needed).
- Ensure **every training script** calls these wrappers for final scoring (no custom MAE/RMSE-only selection without also logging the official metrics).
- Add a simple unit test in `tests/` that:
  - Builds tiny fake `df_actual`, `df_pred`, `df_aux`.
  - Calls `compute_metric1`, `compute_metric2`.
  - Asserts basic properties (e.g., perfect prediction → metric = 0).

### 1.3 Standardize Cross-Validation Strategy

- In `src/validation.py`, confirm there is a single function, e.g. `create_group_kfold_by_brand(bucket_stratified=True)`.
- Ensure **all models** (CatBoost, LGBM, XGB, etc.) use that function to get splits.
- Add a test that verifies:
  - All months of a given brand go to the same fold.
  - Bucket 1 and Bucket 2 are reasonably balanced across folds.

---

## 2. Strong & Stable Single-Model Baselines

### 2.1 Re-run Core Models on a Shared Setup

- For **Scenario 1**, run:

- CatBoost
- LightGBM
- XGBoost
- Linear / Ridge / ElasticNet (if implemented)
- Hybrid physics-ML
- ARIHOW
- For **Scenario 2**, run the same family of models.
- Ensure all runs:
  - Use the same CV split function.
  - Use the same feature config for that scenario.
  - Log official metrics via the centralized metric module.

## 2.2 Build a Model Leaderboard

- Parse `training_all_models.log` and any per-run JSON/CSV metrics to create:
  - `artifacts/model_leaderboard.csv` with columns:
    - `scenario`, `model_name`, `seed`, `cv_metric_s1`, `cv_metric_s2`, `config_path`, `timestamp`.
- Add a small script (e.g. `tools/aggregate_model_scores.py`) that regenerates the leaderboard from logs.
- Sort models by the relevant scenario metric and mark **top N** models for each scenario.

## 2.3 Declare "Hero" Models per Scenario

- In `docs/planning/approach.md`, document:
  - The **best single model** for Scenario 1.
  - The **best single model** for Scenario 2.
- Record:
  - Config file used.
  - Official metrics.
  - Notes on stability (e.g., variance across folds / seeds).

# 3. Time & Bucket Weighting Improvements

## 3.1 Review and Refine Sample Weights

- Open the sample weight section in `run_defaults.yaml` (or relevant config).
- Confirm current behavior:
  - Months 0–5, 6–11, 12–23 have different weights.
  - Bucket 1 rows have higher weight than Bucket 2 (e.g. ×2).
- Propose a **small set** of alternative weighting schemes, such as:
  - Slightly higher weight on months 0–5 (Scenario 1) and 6–11 (both scenarios).
  - Slightly higher weight on Bucket 1 (but not extreme).

## 3.2 Controlled Weighting Experiments

- For each candidate weight scheme:
  - Run the **hero** model for Scenario 1 and Scenario 2 with the modified weights.

- Log official metrics and store configs separately (e.g. `configs/experiments/weights_v1.yaml` ).
  - Compare metrics to the baseline scheme and keep the best-performing version.
  - Update `run_defaults.yaml` only if the improvement is consistent across seeds/folds.
- 

## 4. Targeted Hyperparameter Optimization (HPO)

### 4.1 Define Tight Search Spaces

- For CatBoost, LGBM, and XGBoost configs:
  - Add HPO ranges for:
    - `learning_rate`
    - `depth` / `max_depth`
    - `n_estimators` (if using early stopping, allow higher max)
    - `l2_leaf_reg` / regularization parameters.
  - Keep ranges **narrow, sane, and tabular-TS-friendly**.
- Document the search spaces in comments within `model_*.yaml` .

### 4.2 Enable and Run Small HPO Sweeps

- Ensure `train.py` supports an HPO mode (`--hpo` or similar).
- For each scenario:
  - Run a **small** sweep for CatBoost (and optionally one GBM alternative).
  - Use the **official metric** as the optimization objective.
- Save best hyperparameters in a separate config file (e.g. `configs/model_catboost_s1_best.yaml` ).

### 4.3 Freeze and Document Best Hyperparams

- Update the main configs to use the HPO results (or reference them).
  - Keep the original defaults as commented blocks or separate “baseline” configs.
  - Note in `docs/planning/approach.md` which hyperparameters came from HPO.
- 

## 5. Smarter Ensembles

### 5.1 Prepare Base Models for Ensembling

- Make sure all base models used in ensembles:
  - Share the same CV splits.
  - Use the same target scaling (if any).
  - Produce predictions for the same set of brands and months.
- Add a helper in `src/models/ensemble_utils.py` to align predictions by:
  - `country`, `brand_name`, `months_postgx` keys.

### 5.2 Build and Evaluate Simple Ensembles

- Implement or reuse simple ensemble classes:
  - Plain average (equal weights).

- Weighted average (weights learned on validation).
- Simple stacking (meta-model over base predictions).
- Test ensembles like:
  - CatBoost + LightGBM.
  - CatBoost + Hybrid + ARIHOW.
- For each ensemble:
  - Compute Scenario 1 & Scenario 2 metrics.
  - Save ensemble composition and weights in JSON/YAML under `artifacts/ensembles/`.

## 5.3 Add Guardrails to Ensemble Predictions

- Implement a `sanitize_predictions(df_pred)` function that:
    - Clips negative volumes to 0.
    - Optionally clips very extreme values (e.g.  $> 5 \times$  pre-entry avg) or at least flags them.
  - Call `sanitize_predictions()` before generating submission files.
  - Add a small test to ensure sanitization behaves as expected.
- 

# 6. Inference & Submission Robustness

## 6.1 Scenario Detection & Fallbacks

- Inspect `inference.py` logic for detecting Scenario 1 vs Scenario 2 test series.
- Add tests for:
  - Proper scenario classification given different combinations of available months.
  - Edge cases (e.g. missing some early months).
- Verify fallback strategies (e.g. exponential decay) are triggered when:
  - A model does not produce a prediction for a series/month.
  - The prediction is NaN or obviously invalid.

## 6.2 Submission Sanity Checks

- Implement a small script (`tools/check_submission.py`) that:
    - Reads a submission CSV.
    - Verifies:
      - Correct column names and types.
      - No negative volumes.
      - Reasonable range vs pre-entry `avg_vol` from train data.
    - Prints summary stats per scenario and bucket.
  - Run this script for every candidate submission before uploading.
- 

# 7. Experiment Tracking & Presentation Story

## 7.1 Experiment Tracking

- Ensure each training run:
  - Logs model name, scenario, config path, seed, official metrics to a centralized log (e.g. `training_all_models.log` + CSV).

- Add or improve a small `ExperimentTracker` utility that:
  - Appends entries to a structured file (`experiments.csv` or similar).
  - Can be queried (e.g. via a notebook) to list best runs.

## 7.2 Presentation Outline

- Create `docs/presentation_outline.md` with sections:
    - Problem framing (LOE, generic erosion, scenarios 1 & 2).
    - Data understanding (brands, buckets, erosion behavior).
    - Modeling approach (features, validation, models, ensembles).
    - Results and insights (official metrics, bucket-wise performance).
    - Business interpretation (why early months and Bucket 1 matter).
- 

## 8. How Copilot Should Behave in This Repo

- Favor **small, incremental edits** over large refactors.
  - Suggest moving parameters into config rather than duplicating logic.
  - Auto-complete patterns based on existing code:
    - New model classes should mimic current `BaseModel` subclasses.
    - New tests follow patterns from `tests/test_smoke.py`.
  - Frequently remind the user (via comments or docstrings) to:
    - Re-run `pytest` after non-trivial changes.
    - Re-run a known baseline config to check no performance regression.
    - Update docs and experiment tracker when a new "hero" result is found.
- 

### End goal:

A **clean, well-documented, and reproducible** Arman project with:

- Strong baselines per scenario,
- Thoughtful weighting and HPO,
- Smart ensembles,
- Robust inference/submission code,
- And a clear evidence trail for why your final submissions deserve to top the leaderboard.