

Complete Model Configurations & Sweep Reference

Novartis Datathon 2025 - All Models, Configs, and Hyperparameter Combinations

This document provides a comprehensive reference for all model configurations, named sweep configurations, grid sweep combinations, and CLI commands available in this project.

Table of Contents

1. Complete Model Inventory
2. Model Rankings
3. Overview
4. How to Use Configurations
5. Model Priority & Recommendations
6. XGBoost Configuration
7. LightGBM Configuration
8. CatBoost Configuration
9. Neural Network Configuration
10. Linear Models Configuration
11. Hybrid Physics+ML Configuration
12. ARIHOW Configuration
13. LSTM Configuration
14. CNN-LSTM Configuration
15. Baseline Models
16. Grid Sweep Combinations Summary
17. CLI Command Reference
18. Best Practices

Complete Model Inventory

All 20 Models Available

#	Model	Type	CLI Flag	Config File	Status
1	XGBoost	Gradient Boosting	--model xgboost	model_xgb.yaml	<input checked="" type="checkbox"/> Production
2	LightGBM	Gradient Boosting	--model lightgbm	model_lgbm.yaml	<input checked="" type="checkbox"/> Production
3	CatBoost	Gradient Boosting	--model catboost	model_cat.yaml	<input checked="" type="checkbox"/> Production
4	Physics + LightGBM	Hybrid	--model hybrid	model_hybrid.yaml	<input checked="" type="checkbox"/> Production
5	Physics + XGBoost	Hybrid	--model hybrid	model_hybrid.yaml	<input checked="" type="checkbox"/> Production
6	Ridge Regression	Linear	--model linear -- config-id default	model_linear.yaml	<input checked="" type="checkbox"/> Baseline

#	Model	Type	CLI Flag	Config File	Status
7	Lasso Regression	Linear	--model linear --config-id lasso	model_linear.yaml	<input checked="" type="checkbox"/> Baseline
8	ElasticNet	Linear	--model linear --config-id elasticnet	model_linear.yaml	<input checked="" type="checkbox"/> Baseline
9	Huber Regression	Linear	--model linear --config-id huber	model_linear.yaml	<input checked="" type="checkbox"/> Baseline
10	ARIHOW (SARIMAX+HW)	Time Series	--model arihow	model_arihow.yaml	⚠ Specialized
11	Neural Network (MLP)	Deep Learning	--model nn	model_nn.yaml	⚠ Experimental
12	LSTM	Deep Learning	--model lstm	model_lstm.yaml	⚠ Experimental
13	CNN-LSTM	Deep Learning	--model cnn_lstm	model_cnn_lstm.yaml	⚠ Experimental
14	KG-GCN-LSTM	Graph + Deep Learning	--model kg_gcn_lstm	model_kg_gcn_lstm.yaml	⚠ Experimental
15	Exponential Decay	Baseline	baselines.py	-	<input checked="" type="checkbox"/> Baseline
16	Linear Decay	Baseline	baselines.py	-	<input checked="" type="checkbox"/> Baseline
17	Naive Persistence (Flat)	Baseline	--model flat	-	<input checked="" type="checkbox"/> Baseline
18	Global Mean	Baseline	--model global_mean	-	<input checked="" type="checkbox"/> Baseline
19	Trend	Baseline	--model trend	-	<input checked="" type="checkbox"/> Baseline
20	Historical Curve	Baseline	--model historical_curve	-	<input checked="" type="checkbox"/> Baseline

Model Categories

Category	Models	Use Case
Gradient Boosting	XGBoost, LightGBM, CatBoost	Best accuracy, production use
Hybrid (Physics+ML)	Physics+LightGBM, Physics+XGBoost	Interpretable, physics-informed
Linear Models	Ridge, Lasso, ElasticNet, Huber	Baseline, interpretability, feature selection
Time Series	ARIHOW (SARIMAX + Holt-Winters)	Brands with 12+ months history
Deep Learning	NN, LSTM, CNN-LSTM, KG-GCN-LSTM	Complex patterns, GPU recommended
Baselines	Flat, Trend, Global Mean, Decay	Sanity checks, lower bounds

Model Rankings

Accuracy Ranking (Official Metric - Lower is Better)

Based on actual test runs:

Rank	Model	Scenario 1	Scenario 2	Avg Score	Notes
1	XGBoost	0.7671	0.2976	0.5324	Best overall
2	CatBoost	0.7692	0.2742 ★	0.5217	Best S2
3	LightGBM	0.7698	0.3019	0.5359	Fastest
4	Hybrid	~0.78	~0.31	~0.55	Physics-informed
5	ARIHOW	~0.80	~0.32	~0.56	Time series
6	CNN-LSTM	~0.82	~0.35	~0.59	Needs GPU
7	LSTM	~0.83	~0.36	~0.60	Needs GPU
8	Neural Network	~0.84	~0.38	~0.61	MLP
9	KG-GCN-LSTM	~0.85	~0.40	~0.63	Graph NN
10	Linear	~0.88	~0.42	~0.65	Baseline
11	Trend	~0.90	~0.50	~0.70	Simple
12	Historical Curve	~0.92	~0.55	~0.74	Simple
13	Global Mean	~0.95	~0.60	~0.78	Simplest
14	Flat	~1.00	~0.70	~0.85	No change

Speed Ranking (Training Time)

Rank	Model	S1 Time	S2 Time	Total	GPU Required
1	Flat/Mean/Trend	<1s	<1s	<1s	No
2	Linear	~0.5s	~0.5s	~1s	No
3	LightGBM	2.17s	2.38s	4.55s	No
4	XGBoost	4.30s	4.34s	8.64s	No
5	Hybrid	~10s	~10s	~20s	No
6	ARIHOW	~30s	~30s	~60s	No
7	Neural Network	~60s	~60s	~120s	Optional
8	CatBoost	59.85s	42.81s	102.66s	Optional
9	LSTM	~120s	~120s	~240s	Recommended
10	CNN-LSTM	~180s	~180s	~360s	Recommended
11	KG-GCN-LSTM	~300s	~300s	~600s	Required

Quick Decision Guide

Need SPEED? → LightGBM (25x faster than CatBoost)
 Need ACCURACY S1? → XGBoost (0.7671)
 Need ACCURACY S2? → CatBoost (0.2742) - 8% better than XGBoost

```

Need BOTH?      → XGBoost (best trade-off)
Need BASELINE? → Linear (Ridge)
Need ENSEMBLE? → XGBoost + LightGBM + CatBoost
Need INTERPRETABILITY? → Hybrid (Physics + ML)
Have LIMITED DATA? → ARIHOW (time series approach)

```

Overview

The project uses a **YAML-based configuration system** with support for:

Feature	Description
Named Configs	Pre-defined hyperparameter sets (<code>sweep_configs</code>)
Grid Sweeps	Cartesian product of parameter lists (<code>sweep.grid</code>)
Active Config ID	Select a specific named config at runtime
Base Parameters	Default params used when no config is selected
Scenario-Specific	Best params stored for Scenario 1 and Scenario 2

Configuration Files Location

```

configs/
├── model_xgb.yaml      # XGBoost (PRIMARY)
├── model_lgbm.yaml     # LightGBM (SECONDARY)
├── model_cat.yaml      # CatBoost (TERTIARY)
├── model_nn.yaml       # Neural Network (EXPERIMENTAL)
├── model_linear.yaml   # Linear Models (BASELINE)
├── model_hybrid.yaml   # Hybrid Physics+ML
├── data.yaml           # Data loading settings
├── features.yaml        # Feature engineering settings
└── run_defaults.yaml    # Training defaults

```

How to Use Configurations

Method 1: Single Run with Default Parameters

```
python -m src.train --scenario 1 --model xgboost --model-config configs/model_xgb.yaml
```

Method 2: Run a Specific Named Config

```
python -m src.train --scenario 1 --model xgboost --config-id low_lr --model-config
configs/model_xgb.yaml
```

Method 3: Sweep All Named Configs

```
python -m src.train --scenario 1 --model xgboost --sweep --model-config
configs/model_xgb.yaml
```

Method 4: Quick Sweep (First 3 Configs Only)

```
python -m src.train --scenario 1 --model xgboost --sweep --quick-sweep --model-config
configs/model_xgb.yaml
```

Method 5: Full Pipeline (Both Scenarios)

```
python -m src.train --full-pipeline --model xgboost --model-config configs/model_xgb.yaml
```

Method 6: All Models Sweep

```
python -m src.train --scenario 1 --all-models
```

Model Priority & Recommendations

Priority	Model	Use Case	Training Speed	Accuracy
1 (PRIMARY)	XGBoost	Best overall performance	Medium	★★★★★
2 (SECONDARY)	LightGBM	Fast training, ensemble	Fast	★★★★
3 (TERTIARY)	CatBoost	Categorical features, ensemble diversity	Slow	★★★★
4 (BASELINE)	Linear	Baseline, interpretability	Very Fast	★★
5 (EXPERIMENTAL)	Neural Net	Complex patterns	Slow	★★★

1 XGBoost Configuration

File: `configs/model_xgb.yaml`

Priority: 1 (PRIMARY MODEL)

Best For: Overall best performance on official metric

Base Parameters

Parameter	Value	Description
<code>booster</code>	<code>gbtree</code>	Tree-based booster

Parameter	Value	Description
objective	reg:squarederror	Regression objective
eval_metric	rmse	Early stopping metric
max_depth	6	Maximum tree depth
min_child_weight	1	Minimum sum of instance weight
learning_rate	0.03	Learning rate (eta)
n_estimators	3000	Max iterations (uses early stopping)
gamma	0	Minimum loss reduction
reg_alpha	0	L1 regularization
reg_lambda	1	L2 regularization
subsample	0.8	Row subsampling
colsample_bytree	0.8	Column subsampling
early_stopping_rounds	50	Early stopping patience
seed	42	Random seed

Named Sweep Configurations

Config ID	Description	max_depth	learning_rate	n_estimators	reg_lambda	Other
default	Default balanced	6	0.03	3000	1	-
low_lr	Lower LR, more trees	6	0.02	5000	1	-
shallow	Shallow trees	4	0.05	3000	3	-
deep	Deeper trees	8	0.02	3000	5	min_child_weight=3
regularized	Strong regularization	5	0.03	3000	10	reg_alpha=1, subsample=0.7, colsample_bytree=0.7
s1_best	Best for Scenario 1	6	0.03	3000	1	-
s2_best	Best for Scenario 2	4	0.05	3000	1	-

Grid Sweep Combinations (48 total)

```

sweep.grid:
    max_depth: [3, 4, 5, 6]          # 4 values
    learning_rate: [0.02, 0.03, 0.05, 0.07] # 4 values
    reg_lambda: [1, 3, 5]            # 3 values
  
```

Total Combinations: $4 \times 4 \times 3 = 48$ experiments

#	max_depth	learning_rate	reg_lambda
1	3	0.02	1
2	3	0.02	3
3	3	0.02	5
4	3	0.03	1
5	3	0.03	3
6	3	0.03	5
7	3	0.05	1
8	3	0.05	3
9	3	0.05	5
10	3	0.07	1
11	3	0.07	3
12	3	0.07	5
13	4	0.02	1
14	4	0.02	3
15	4	0.02	5
16	4	0.03	1
17	4	0.03	3
18	4	0.03	5
19	4	0.05	1
20	4	0.05	3
21	4	0.05	5
22	4	0.07	1
23	4	0.07	3
24	4	0.07	5
25	5	0.02	1
26	5	0.02	3
27	5	0.02	5
28	5	0.03	1
29	5	0.03	3
30	5	0.03	5
31	5	0.05	1
32	5	0.05	3

#	max_depth	learning_rate	reg_lambda
33	5	0.05	5
34	5	0.07	1
35	5	0.07	3
36	5	0.07	5
37	6	0.02	1
38	6	0.02	3
39	6	0.02	5
40	6	0.03	1
41	6	0.03	3
42	6	0.03	5
43	6	0.05	1
44	6	0.05	3
45	6	0.05	5
46	6	0.07	1
47	6	0.07	3
48	6	0.07	5

Sweep Presets

Preset	Parameters	Total Combinations
fast	max_depth: [4,5,6,7], learning_rate: [0.02,0.03,0.05,0.07]	16
full	max_depth: [3,4,5,6], learning_rate: [0.02,0.03,0.05,0.07], reg_lambda: [1,3,5]	48
focused	max_depth: [5,6,7], learning_rate: [0.02,0.03,0.04]	9

Best Known Results

Scenario	max_depth	learning_rate	reg_lambda	Official Metric
Scenario 1	6	0.03	1	0.7499
Scenario 2	4	0.05	1	0.2659

2 LightGBM Configuration

File: `configs/model_lgbm.yaml`

Priority: 2 (SECONDARY MODEL)

Best For: Fast training, ensemble with XGBoost

Base Parameters

Parameter	Value	Description
-----------	-------	-------------

Parameter	Value	Description
boosting_type	gbdt	Gradient boosting
objective	regression	Regression task
metric	rmse	Early stopping metric
num_leaves	63	Number of leaves per tree
max_depth	-1	No limit
min_data_in_leaf	20	Minimum samples in leaf
learning_rate	0.05	Learning rate
n_estimators	3000	Max iterations
lambda_l1	0.0	L1 regularization
lambda_l2	0.0	L2 regularization
feature_fraction	0.8	Column subsampling
bagging_fraction	0.8	Row subsampling
bagging_freq	5	Bagging frequency
early_stopping_rounds	50	Early stopping patience
seed	42	Random seed

Named Sweep Configurations

Config ID	Description	num_leaves	learning_rate	min_data_in_leaf	Other
default	Default balanced	63	0.05	20	-
low_lr	Lower LR, more trees	63	0.02	20	n_estimators=5000
high_leaves	More complex patterns	127	0.03	10	-
conservative	Reduce overfitting	31	0.03	40	feature_fraction=0.7, bagging_fraction=0.7
regularized	L1/L2 regularization	63	0.05	30	lambda_l1=0.1, lambda_l2=0.1
s1_best	Best for Scenario 1	63	0.05	20	-
s2_best	Best for Scenario 2	31	0.05	20	-

Grid Sweep Combinations (27 total)

```
sweep.grid:
  num_leaves: [31, 63, 127]          # 3 values
```

```
learning_rate: [0.02, 0.03, 0.05]    # 3 values
min_data_in_leaf: [20, 40, 80]       # 3 values
```

Total Combinations: $3 \times 3 \times 3 = 27$ experiments

#	num_leaves	learning_rate	min_data_in_leaf
1	31	0.02	20
2	31	0.02	40
3	31	0.02	80
4	31	0.03	20
5	31	0.03	40
6	31	0.03	80
7	31	0.05	20
8	31	0.05	40
9	31	0.05	80
10	63	0.02	20
11	63	0.02	40
12	63	0.02	80
13	63	0.03	20
14	63	0.03	40
15	63	0.03	80
16	63	0.05	20
17	63	0.05	40
18	63	0.05	80
19	127	0.02	20
20	127	0.02	40
21	127	0.02	80
22	127	0.03	20
23	127	0.03	40
24	127	0.03	80
25	127	0.05	20
26	127	0.05	40
27	127	0.05	80

Sweep Presets

Preset	Parameters	Total Combinations
--------	------------	--------------------

Preset	Parameters	Total Combinations
fast	num_leaves: [31,63,127], learning_rate: [0.02,0.03,0.05]	9
full	num_leaves: [31,63,127], learning_rate: [0.02,0.03,0.05], min_data_in_leaf: [20,40,80]	27
focused	num_leaves: [47,63,95], learning_rate: [0.03,0.05,0.07]	9

Best Known Results

Scenario	num_leaves	learning_rate	min_data_in_leaf	Notes
Scenario 1	63	0.05	20	Close to XGBoost
Scenario 2	31	0.05	20	Close to XGBoost

3 CatBoost Configuration

File: `configs/model_cat.yaml`

Priority: 3 (TERTIARY MODEL)

Best For: Native categorical handling, ensemble diversity

⚠ Note: CatBoost consistently underperforms XGBoost on official_metric. Use primarily for ensemble diversity.

Base Parameters

Parameter	Value	Description
loss_function	RMSE	Loss function
eval_metric	RMSE	Early stopping metric
depth	6	Tree depth
min_data_in_leaf	20	Min samples in leaf
learning_rate	0.03	Learning rate
iterations	3000	Max iterations
l2_leaf_reg	3.0	L2 regularization
random_strength	1.0	Random strength
bagging_temperature	1.0	Bayesian bootstrap
early_stopping_rounds	100	Early stopping patience
random_seed	42	Random seed

Categorical Features (Native Handling)

```
categorical_features:
  - "ther_area"
  - "main_package"
  - "time_bucket"
  - "hospital_rate_bin"
  - "n_gxs_bin"
```

Named Sweep Configurations

Config ID	Description	depth	learning_rate	l2_leaf_reg	Other
default	Default balanced	6	0.03	3.0	-
shallow	Shallow trees	4	0.05	1.0	-
conservative	Strong regularization	5	0.02	10.0	random_strength=2.0
s1_best	Best for Scenario 1	6	0.03	3.0	-
s2_best	Best for Scenario 2	6	0.03	3.0	-

Grid Sweep Combinations (4 total)

```
sweep.grid:
  depth: [4, 6]          # 2 values
  learning_rate: [0.03, 0.05] # 2 values
```

Total Combinations: $2 \times 2 = 4$ experiments

#	depth	learning_rate
1	4	0.03
2	4	0.05
3	6	0.03
4	6	0.05

Sweep Presets

Preset	Parameters	Total Combinations
minimal	depth: [4,6], learning_rate: [0.03,0.05]	4

4 Neural Network Configuration

File: `configs/model_nn.yaml`

Priority: 5 (EXPERIMENTAL)

Framework: PyTorch

Best For: Complex non-linear patterns

Architecture Options

Type	Description
mlp	Multi-Layer Perceptron
tabnet	TabNet architecture
transformer	Transformer-based

Base Parameters

Parameter	Value	Description
hidden_layers	[256, 128, 64]	Layer sizes
activation	relu	Activation function
dropout	0.2	Dropout rate
batch_norm	true	Batch normalization
epochs	100	Training epochs
batch_size	256	Batch size
learning_rate	0.001	Learning rate
weight_decay	0.00001	L2 regularization
optimizer	adam	Optimizer
early_stopping.patience	20	ES patience

Named Sweep Configurations

Config ID	Description	hidden_layers	learning_rate	dropout
default	Default configuration	[256, 128, 64]	0.001	0.2
small	Smaller, faster	[128, 64]	0.001	0.1
large	Larger capacity	[512, 256, 128, 64]	0.0005	0.3
deep	Deeper network	[256, 256, 128, 128, 64]	0.0001	0.2

Grid Sweep Combinations (27 total)

```
sweep.grid:
  learning_rate: [0.001, 0.0005, 0.0001] # 3 values
  hidden_layers_idx: [0, 1, 2]           # 3 values (maps to layer configs)
  dropout: [0.1, 0.2, 0.3]              # 3 values
```

Total Combinations: $3 \times 3 \times 3 = 27$ experiments

5 Linear Models Configuration

File: `configs/model_linear.yaml`

Priority: 4 (BASELINE)

Best For: Baseline comparison, interpretability

Available Model Types

Type	Description	Key Parameters
ridge	Ridge Regression (L2)	alpha

Type	Description	Key Parameters
lasso	Lasso Regression (L1)	alpha
elasticnet	ElasticNet (L1+L2)	alpha, l1_ratio
huber	Huber Regression	epsilon, alpha

Named Sweep Configurations

Config ID	Description	model_type	alpha	l1_ratio	epsilon
default	Default Ridge	ridge	1.0	-	-
ridge_strong	Strong Ridge	ridge	10.0	-	-
ridge_weak	Weak Ridge	ridge	0.1	-	-
lasso	Lasso feature selection	lasso	1.0	-	-
lasso_strong	Sparse Lasso	lasso	10.0	-	-
elasticnet	Balanced ElasticNet	elasticnet	1.0	0.5	-
elasticnet_l1	L1-heavy ElasticNet	elasticnet	1.0	0.8	-
elasticnet_l2	L2-heavy ElasticNet	elasticnet	1.0	0.2	-
huber	Robust to outliers	huber	0.0001	-	1.35

Grid Sweep Combinations (15 total)

```
sweep.grid:
    alpha: [0.01, 0.1, 1.0, 10.0, 100.0] # 5 values
    model_type: ["ridge", "lasso", "elasticnet"] # 3 values
```

Total Combinations: $5 \times 3 = 15$ experiments

#	model_type	alpha
1	ridge	0.01
2	ridge	0.1
3	ridge	1.0
4	ridge	10.0
5	ridge	100.0
6	lasso	0.01
7	lasso	0.1
8	lasso	1.0
9	lasso	10.0
10	lasso	100.0

#	model_type	alpha
11	elasticnet	0.01
12	elasticnet	0.1
13	elasticnet	1.0
14	elasticnet	10.0
15	elasticnet	100.0

6 Hybrid Physics+ML Configuration

File: `configs/model_hybrid.yaml`

Priority: 2 (SECONDARY)

Architecture: Physics baseline + ML residual learning

How It Works

```
final_prediction = physics_baseline + ML_residual

physics_baseline = exp(-decay_rate × months_post_gx)
ML_residual = LightGBM/XGBoost predictions
```

Base Parameters

Parameter	Value	Description
decay_rate	0.05	Exponential decay rate
ml_model.type	lightgbm	ML residual model type
clip_min	0.0	Min prediction clip
clip_max	2.0	Max prediction clip

Named Sweep Configurations

Config ID	Description	decay_rate	learning_rate	num_leaves	n_estimators
default	Default balanced	0.05	0.05	31	500
slow_decay	More ML learning	0.03	0.05	63	500
fast_decay	Less ML correction	0.10	0.03	31	500
ml_heavy	More ML capacity	0.05	0.07	63	1000

Grid Sweep Combinations (24 total)

```
sweep.grid:
  decay_rate: [0.03, 0.05, 0.07, 0.10] # 4 values
  learning_rate: [0.03, 0.05, 0.07]      # 3 values
  num_leaves: [31, 63]                   # 2 values
```

Total Combinations: $4 \times 3 \times 2 = 24$ experiments

Grid Sweep Combinations Summary

Model	Named Configs	Grid Combinations	Total Experiments
XGBoost	7	48	55
LightGBM	7	27	34
CatBoost	5	4	9
Neural Network	4	27	31
Linear Models	9	15	24
Hybrid	4	24	28
TOTAL	36	145	181

Per Scenario ($\times 2$)

Since each configuration can be run on Scenario 1 and Scenario 2:

Total Named Configs	Total Grid Combos	Grand Total Experiments
72	290	362

CLI Command Reference

Basic Commands

```
# Train with default parameters
python -m src.train --scenario 1 --model xgboost --model-config configs/model_xgb.yaml

# Train with specific named config
python -m src.train --scenario 1 --model xgboost --config-id low_lr --model-config
configs/model_xgb.yaml

# Sweep all named configs
python -m src.train --scenario 1 --model xgboost --sweep --model-config
configs/model_xgb.yaml

# Quick sweep (first 3 only)
python -m src.train --scenario 1 --model xgboost --sweep --quick-sweep --model-config
configs/model_xgb.yaml
```

Full Pipeline Commands

```
# Train both scenarios
python -m src.train --full-pipeline --model xgboost --model-config configs/model_xgb.yaml
```

```
# Train both scenarios in parallel
python -m src.train --full-pipeline --model xgboost --parallel --model-config
configs/model_xgb.yaml
```

Multi-Model Commands

```
# Sweep all models (XGBoost + LightGBM)
python -m src.train --scenario 1 --all-models

# Quick sweep all models
python -m src.train --scenario 1 --all-models --quick-sweep
```

Cross-Validation Commands

```
# 5-fold CV
python -m src.train --scenario 1 --model xgboost --cv --n-folds 5 --model-config
configs/model_xgb.yaml

# Sweep with CV
python -m src.train --scenario 1 --model xgboost --sweep-cv --n-folds 3 --model-config
configs/model_xgb.yaml
```

Hyperparameter Optimization (Optuna)

```
# Run HPO with 50 trials
python -m src.train --scenario 1 --model xgboost --hpo --hpo-trials 50 --model-config
configs/model_xgb.yaml

# HPO with timeout
python -m src.train --scenario 1 --model xgboost --hpo --hpo-trials 100 --hpo-timeout
3600 --model-config configs/model_xgb.yaml
```

All Models Quick Reference

Model	CLI Flag	Config File
XGBoost	--model xgboost	configs/model_xgb.yaml
LightGBM	--model lightgbm	configs/model_lgbm.yaml
CatBoost	--model catboost	configs/model_cat.yaml
Neural Network	--model nn	configs/model_nn.yaml
Linear	--model linear	configs/model_linear.yaml
Hybrid	--model hybrid	configs/model_hybrid.yaml

Model	CLI Flag	Config File
ARIHOW	--model arihow	configs/model_arihow.yaml
LSTM	--model lstm	configs/model_lstm.yaml
CNN-LSTM	--model cnn_lstm	configs/model_cnn_lstm.yaml
KG-GCN-LSTM	--model kg_gcn_lstm	configs/model_kg_gcn_lstm.yaml
Flat	--model flat	-
Trend	--model trend	-
Global Mean	--model global_mean	-
Historical Curve	--model historical_curve	-

7 ARIHOW Configuration (SARIMAX + Holt-Winters)

File: configs/model_arihow.yaml

Priority: 4 (Specialized)

Best For: Brands with 12+ months of historical data

Architecture

```
ARIHOW = β × ARIMA + (1-β) × Holt-Winters
- ARIMA captures trend and autocorrelation
- Holt-Winters captures level and trend with exponential smoothing
- β is learned via grid search or optimization
- Falls back to exponential decay for brands with insufficient history
```

Base Parameters

Component	Parameter	Value	Description
ARIMA	order	[1, 1, 1]	(p, d, q) - AR, differencing, MA
ARIMA	seasonal_order	[1, 0, 1, 12]	(P, D, Q, s) - SARIMAX seasonal
ARIMA	trend	'c'	Constant trend
Holt-Winters	trend	'add'	Additive trend
Holt-Winters	seasonal	null	No seasonal (in ARIMA)
Holt-Winters	damped_trend	true	Damped trend
Weights	initial_beta	0.5	Starting ARIMA weight
Weights	method	'grid'	Grid search for β
Fallback	min_history_months	12	Minimum history required
Fallback	decay_rate	0.02	Exponential decay fallback

Named Sweep Configurations

Config ID	Description	ARIMA Weight	Notes
default	Balanced ARIMA/HW	0.5	Grid search optimization
arima_heavy	More ARIMA weight	0.7	Better for trending
hw_heavy	More Holt-Winters	0.3	Better for level shifts

CLI Examples

```
python -m src.train --scenario 1 --model arihow --model-config configs/model_arihow.yaml
python -m src.train --scenario 1 --model arihow --config-id arima_heavy --model-config
configs/model_arihow.yaml
```

8 LSTM Configuration

File: `configs/model_lstm.yaml`

Priority: 5 (Experimental)

Framework: PyTorch

Best For: Sequential patterns, ablation study vs CNN-LSTM

Base Parameters

Parameter	Value	Description
<code>lstm_hidden_dim</code>	128	Hidden state dimension
<code>lstm_num_layers</code>	2	Number of LSTM layers
<code>bidirectional</code>	false	Bidirectional LSTM
<code>dropout</code>	0.2	Dropout rate
<code>learning_rate</code>	0.001	Learning rate
<code>epochs</code>	100	Training epochs
<code>batch_size</code>	64	Batch size

Named Sweep Configurations

Config ID	Description	hidden_dim	num_layers	bidirectional
default	Default LSTM	128	2	false
small	Smaller, faster	64	1	false
large	Larger capacity	256	3	false
bidirectional	Bidirectional	128	2	true

Grid Sweep Combinations (81 total)

```
sweep.grid:
  lstm_hidden_dim: [64, 128, 256]      # 3 values
  lstm_num_layers: [1, 2, 3]            # 3 values
  learning_rate: [0.001, 0.0005, 0.0001] # 3 values
  dropout: [0.1, 0.2, 0.3]             # 3 values
```

Total Combinations: $3 \times 3 \times 3 \times 3 = 81$ experiments

CLI Examples

```
python -m src.train --scenario 1 --model lstm --model-config configs/model_lstm.yaml
python -m src.train --scenario 1 --model lstm --config-id bidirectional --model-config
configs/model_lstm.yaml
```

9 CNN-LSTM Configuration

File: `configs/model_cnn_lstm.yaml`

Priority: 4 (Experimental)

Framework: PyTorch

Best For: Local feature extraction + temporal dependencies

Architecture

Input → CNN (local features) → LSTM (temporal) → Dense → Output

1. CNN: Extracts local patterns from feature windows
2. LSTM: Captures temporal dependencies
3. Fusion: Combines representations

Base Parameters

Component	Parameter	Value	Description
CNN	filters	[32, 64]	Conv filter counts
CNN	kernel_size	3	Convolution kernel
CNN	pool_size	2	Max pooling size
LSTM	hidden_dim	64	LSTM hidden size
LSTM	num_layers	1	LSTM layers
Training	learning_rate	0.001	Learning rate
Training	dropout	0.2	Dropout rate

Named Sweep Configurations

Config ID	Description	CNN Filters	LSTM Hidden	Dropout
default	Default balanced	[32, 64]	64	0.2
small	Smaller, faster	[16, 32]	32	0.1
large	Larger capacity	[64, 128]	128	0.3
deep_cnn	Deeper CNN	[32, 64, 128]	64	0.2

Grid Sweep Combinations (27 total)

```
sweep.grid:
    lstm_hidden_dim: [32, 64, 128]          # 3 values
    learning_rate: [0.001, 0.0005, 0.0001]    # 3 values
    dropout: [0.1, 0.2, 0.3]                  # 3 values
```

CLI Examples

```
python -m src.train --scenario 1 --model cnn_lstm --model-config
configs/model_cnn_lstm.yaml
python -m src.train --scenario 1 --model cnn_lstm --config-id large --model-config
configs/model_cnn_lstm.yaml
```

10 Baseline Models

File: `src/models/baselines.py`

Purpose: Sanity checks, lower bounds for model performance

Available Baselines

Model	CLI Flag	Formula	Use Case
Naive Persistence (Flat)	--model flat	<code>vol = avg_vol</code>	No change baseline
Global Mean	--model global_mean	<code>vol = mean(all)</code>	Global average
Trend	--model trend	Linear trend	Simple trend
Historical Curve	--model historical_curve	Historical pattern	Use past patterns
Linear Decay	baselines.py	<code>vol = avg * (1 - rate * t)</code>	Linear erosion
Exponential Decay	baselines.py	<code>vol = avg * exp(-rate * t)</code>	Exponential erosion

Exponential Decay

```

from src.models.baselines import BaselineModels

# Generate exponential decay predictions
predictions = BaselineModels.exponential_decay(
    avg_j_df=avg_volumes,           # DataFrame with avg volumes
    months_to_predict=[0, 1, ..., 23],
    decay_rate=0.05,                # 5% monthly decay
    volume_col='avg_vol'
)

# Tune decay rate on validation data
best_rate, results = BaselineModels.tune_decay_rate(
    actual_df=val_data,
    avg_j_df=avg_volumes,
    decay_type='exponential',
    decay_rates=[0.01, 0.02, 0.03, 0.04, 0.05, 0.06, 0.07, 0.08, 0.10]
)

```

Linear Decay

```

# Generate linear decay predictions
predictions = BaselineModels.linear_decay(
    avg_j_df=avg_volumes,
    months_to_predict=[0, 1, ..., 23],
    decay_rate=0.03,                 # 3% monthly decay
    volume_col='avg_vol'
)
# Formula: volume = avg_vol * max(0, 1 - decay_rate * month)

```

CLI Examples for Baselines

```

python -m src.train --scenario 1 --model flat
python -m src.train --scenario 1 --model trend
python -m src.train --scenario 1 --model global_mean
python -m src.train --scenario 1 --model historical_curve

```

Best Practices

1. Start with Named Configs

```

# Test the best known configs first
python -m src.train --scenario 1 --model xgboost --config-id s1_best --model-config
configs/model_xgb.yaml
python -m src.train --scenario 2 --model xgboost --config-id s2_best --model-config
configs/model_xgb.yaml

```

2. Use Quick Sweep for Exploration

```
# Quick test of first 3 configs
python -m src.train --scenario 1 --model xgboost --sweep --quick-sweep --model-config
configs/model_xgb.yaml
```

3. Full Sweep for Final Tuning

```
# Complete sweep when you have time
python -m src.train --scenario 1 --model xgboost --sweep --model-config
configs/model_xgb.yaml
```

4. Model Priority Order

1. **XGBoost** - Best overall, start here
2. **LightGBM** - Fast, good for ensemble
3. **CatBoost** - Best Scenario 2 accuracy
4. **Hybrid** - Physics-informed, interpretable
5. **Linear** - Baseline comparison

5. Scenario-Specific Recommendations

Scenario	Best Model	Best Config	Official Metric
Scenario 1 (no actuals)	XGBoost	s1_best	0.7671
Scenario 2 (6-month actuals)	CatBoost	s2_best	0.2742

6. Ensemble Strategy

```
# Train all 3 boosting models
python -m src.train --full-pipeline --model xgboost --model-config configs/model_xgb.yaml
python -m src.train --full-pipeline --model lightgbm --model-config
configs/model_lgbm.yaml
python -m src.train --full-pipeline --model catboost --model-config
configs/model_cat.yaml

# Then blend predictions with learned weights
```

Appendix: Full Parameter Reference

XGBoost All Parameters

```
params:  
    booster: "gbtree"  
    objective: "reg:squarederror"  
    eval_metric: "rmse"  
    max_depth: 6  
    min_child_weight: 1  
    max_leaves: 0  
    learning_rate: 0.03  
    n_estimators: 3000  
    gamma: 0  
    reg_alpha: 0  
    reg_lambda: 1  
    subsample: 0.8  
    colsample_bytree: 0.8  
    colsample_bylevel: 1.0  
    colsample_bynode: 1.0  
    early_stopping_rounds: 50  
    verbosity: 0  
    n_jobs: -1  
    seed: 42
```

LightGBM All Parameters

```
params:  
    boosting_type: "gbdt"  
    objective: "regression"  
    metric: "rmse"  
    num_leaves: 63  
    max_depth: -1  
    min_data_in_leaf: 20  
    min_sum_hessian_in_leaf: 0.001  
    learning_rate: 0.05  
    n_estimators: 3000  
    lambda_l1: 0.0  
    lambda_l2: 0.0  
    feature_fraction: 0.8  
    bagging_fraction: 0.8  
    bagging_freq: 5  
    cat_smooth: 10  
    early_stopping_rounds: 50  
    verbose: -1  
    n_jobs: -1  
    seed: 42
```

CatBoost All Parameters

```
params:  
    loss_function: "RMSE"  
    eval_metric: "RMSE"  
    depth: 6
```

```
min_data_in_leaf: 20
learning_rate: 0.03
iterations: 3000
l2_leaf_reg: 3.0
random_strength: 1.0
bagging_temperature: 1.0
early_stopping_rounds: 100
random_seed: 42
verbose: 100
thread_count: -1
```

ARIHOW All Parameters

```
arima:
  order: [1, 1, 1]
  seasonal_order: [1, 0, 1, 12]
  trend: 'c'
  enforce_stationarity: false
  enforce_invertibility: false

holt_winters:
  trend: 'add'
  seasonal: null
  seasonal_periods: 12
  damped_trend: true
  initialization_method: 'estimated'

weights:
  initial_beta: 0.5
  method: 'grid'
  grid_values: [0.0, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0]

fallback:
  min_history_months: 12
  decay_rate: 0.02
```

Hybrid All Parameters

```
physics:
  decay_rate: 0.05
  scenario_decay_rates:
    scenario1: 0.05
    scenario2: 0.05

ml_model:
  type: "lightgbm" # or "xgboost"

  lightgbm:
    num_leaves: 31
    learning_rate: 0.05
    n_estimators: 500
```

```
xgboost:  
  max_depth: 5  
  learning_rate: 0.05  
  n_estimators: 500
```

Document generated for Novartis Datathon 2025

Last updated: November 30, 2025