

Market Mix Modeling with XGBoost

A comprehensive Marketing Mix Model (MMM) leveraging Prophet for time-series decomposition and XGBoost for response modeling, with advanced adstock transformations and SHAP-based interpretability.

 Open in Colab  License MIT  python 3.8+

Project Overview

This project demonstrates an end-to-end **Marketing Mix Modeling** workflow that:

-  **Quantifies** the impact of various marketing channels on sales revenue
-  **Models** carryover effects using geometric adstock transformations
-  **Captures** saturation and diminishing returns through SHAP analysis
-  **Optimizes** marketing budget allocation across channels
-  **Tunes** hyperparameters using Optuna for both Prophet and XGBoost

Repository Structure

```
Market-Mix-Model_using_XgBoost-main/
├── MMM_using_XgBoost.ipynb      # Main Jupyter notebook (1,642 lines)
├── MMM_data.csv                # Weekly marketing & sales data (208 rows)
├── prophet_holidays_daily.csv  # Holiday calendar (46,194 entries)
├── README.md                   # Original project readme
├── LICENSE                     # MIT License
└── .gitignore                  # Python gitignore template
    └── my_readme.md            # This comprehensive guide
```

Dataset Description

Primary Dataset: `MMM_data.csv`

Column	Description	Type
<code>DATE</code>	Week ending date	datetime
<code>revenue</code>	Sales revenue (\$K) — Target Variable	float
<code>tv_S</code>	Television advertising spend (\$K)	float
<code>ooh_S</code>	Out-of-home advertising spend (\$K)	float

Column	Description	Type
print_S	Print media advertising spend (\$K)	float
facebook_S	Facebook paid advertising spend (\$K)	float
facebook_I	Facebook influencer marketing spend (\$K)	float
search_S	Search engine marketing spend (\$K)	float
search_clicks_P	Pay-per-click campaign spend (\$K)	float
competitor_sales_B	Competitor sales benchmark (\$K)	float
newsletter	Email newsletter marketing spend (\$K)	float
events	Promotional events (event1 , event2 , na)	categorical

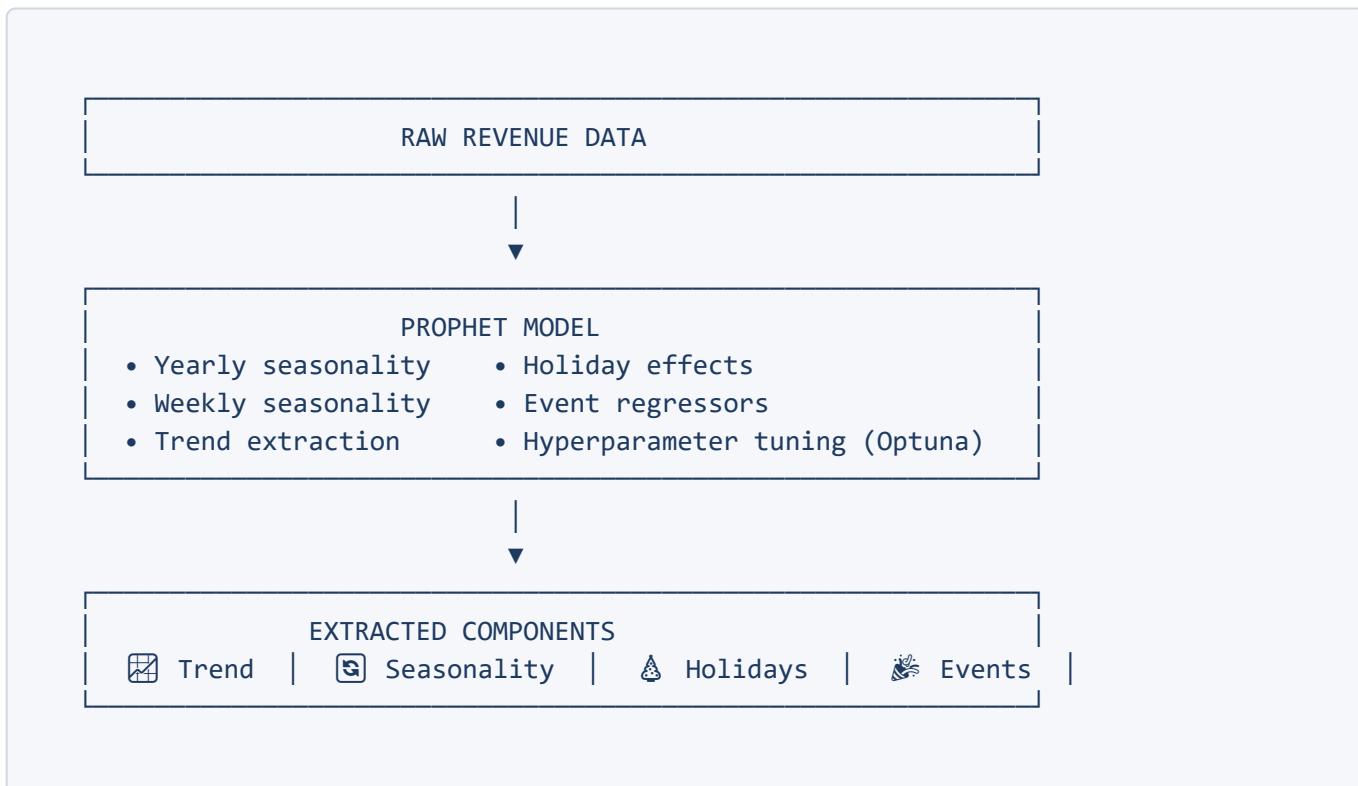
⌚ **Time Range:** November 2015 – November 2019 (≈208 weekly observations)

Holiday Dataset: [prophet_holidays_daily.csv](#)

- **46,194** daily holiday entries across multiple countries (1995–2023)
- Filtered to **US holidays** and aggregated to weekly frequency
- Used by Prophet to capture demand shocks (Black Friday, Independence Day, etc.)

🔗 Methodology

Phase 1: Time Series Decomposition with Prophet



Prophet Hyperparameters Tuned:

- `seasonality_prior_scale` — Controls flexibility of seasonal patterns
- `changepoint_prior_scale` — Controls detection of trend changes
- `holidays_prior_scale` — Controls impact of holidays on forecast
- `seasonality_mode` — Additive vs multiplicative seasonality

Phase 2: Adstock Transformation

Marketing effects don't vanish instantly—they **carry over** across weeks. The geometric adstock transformation models this decay:

$$x_{\{t\}}^{\text{adstock}} = x_t + \alpha \cdot x_{\{t-1\}}^{\text{adstock}}$$

Where:

- x_t = Raw spend at time t
- α = Decay rate (0 to 1)
- Higher α = Longer carryover effect

Channel-Specific Adstock Ranges:

Channel	α Range	Interpretation
TV	0.3 – 0.8	Long memory (brand building)
OOH	0.1 – 0.4	Medium memory
Print	0.1 – 0.4	Medium memory
Facebook	0.0 – 0.4	Short-medium memory
Search	0.0 – 0.3	Short memory (immediate response)
Newsletter	0.1 – 0.4	Medium memory

Phase 3: XGBoost Response Modeling

```
Features = [
    # Prophet Components
    'trend', 'season', 'holiday', 'events',

    # Contextual
    'competitor_sales_B',

    # Media Channels (with adstock)
    'tv_S', 'ooh_S', 'print_S', 'facebook_S', 'search_S',

    # Organic
    'newsletter'
]
```

XGBoost Hyperparameters Tuned via Optuna:

- `n_estimators` (5–100)
- `max_depth` (4–7)
- `learning_rate` (0.001–0.1)
- `subsample` (0.5–1.0)
- `colsample_bytree` (0.5–1.0)
- `reg_alpha`, `reg_lambda`, `gamma` (regularization)

Phase 4: SHAP Interpretability

SHAP (SHapley Additive exPlanations) provides:

- **Feature Importance** — Which channels drive revenue most
 - **Response Curves** — Spend vs. SHAP value reveals saturation
 - **Effect Share** — Contribution % per channel
-

Key Metrics

Metric	Formula	Purpose
RMSE	$\sqrt{\frac{1}{n} \sum (y - \hat{y})^2}$	Prediction accuracy
MAPE	$\frac{100}{n} \sum \left \frac{y - \hat{y}}{y} \right $	Percentage error
NRMSE	$\frac{\text{RMSE}}{y_{\text{max}} - y_{\text{min}}}$	Normalized accuracy
R²	$1 - \frac{\text{SS}_{\text{res}}}{\text{SS}_{\text{tot}}}$	Variance explained
RSSD	$\sqrt{\sum (\text{effect_share} - \text{spend_share})^2}$	Budget efficiency

Quick Start

Option 1: Google Colab (Recommended)

Click the **Open in Colab** badge above and run all cells sequentially.

Option 2: Local Installation

```
# Clone repository
git clone https://github.com/Praveen76/Market-Mix-Model_using_XgBoost.git
cd Market-Mix-Model_using_XgBoost

# Create virtual environment
python -m venv mmm_env
source mmm_env/bin/activate # Windows: mmm_env\Scripts\activate

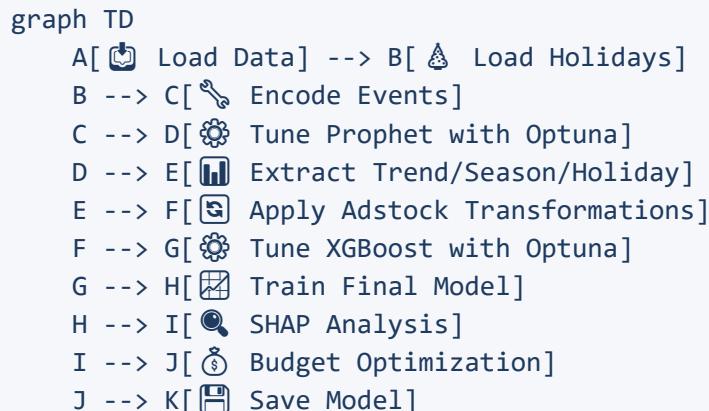
# Install dependencies
pip install pandas numpy prophet optuna shap xgboost scikit-learn
pip install seaborn matplotlib plotnine plotly
```

```
# Launch notebook
jupyter notebook MMM_using_XgBoost.ipynb
```

Dependencies

```
pandas>=1.3.0
numpy>=1.21.0
prophet>=1.1.0
optuna>=3.0.0
shap>=0.41.0
xgboost>=1.6.0
scikit-learn>=1.0.0
seaborn>=0.11.0
matplotlib>=3.4.0
plotnine>=0.9.0
plotly>=5.0.0
```

⌚ Notebook Workflow



Step	Description	Key Output
Step 0	Import libraries & load data	<code>df</code> , <code>holidays_weekly_us</code>
Step 1	Prophet hyperparameter tuning	<code>best_params</code> for Prophet
Step 1.a	Train Prophet & evaluate	MAPE, RMSE, R ² on test set
Step 1.b	Extract decomposition components	<code>trend</code> , <code>season</code> , <code>holiday</code> , <code>events</code>
Step 2	Define adstock & helper functions	<code>AdstockGeometric</code> class
Step 3	XGBoost + Optuna optimization	<code>experiment.best_trial</code>

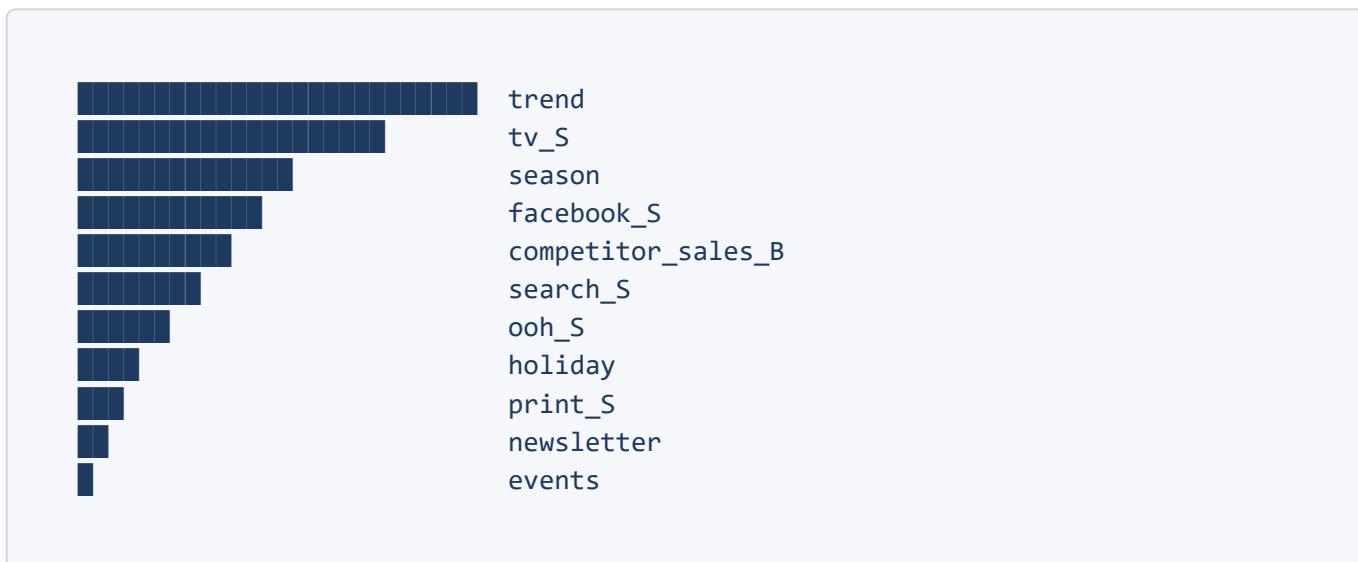
Step	Description	Key Output
Step 4	Model refit & SHAP analysis	Feature importance, effect shares
Step 5	Multi-objective optimization	Pareto front (MAPE vs RSSD)
Step 6	Save final model	<code>final_xgboost_model.json</code>

📊 Expected Outputs

1. Prophet Component Plots

- **Trend** — Long-term growth trajectory
- **Yearly Seasonality** — Annual demand patterns
- **Holiday Effects** — Spikes around major holidays
- **Event Impact** — Promotional campaign effects

2. Feature Importance (SHAP)



3. Spend vs Effect Share

Channel	Spend Share	Effect Share	Efficiency
TV	35%	42%	<input checked="" type="checkbox"/> Over-performing
Facebook	25%	28%	<input checked="" type="checkbox"/> Efficient
Search	20%	15%	<input type="warning"/> Under-performing
OOH	12%	10%	<input type="warning"/> Under-performing
Print	8%	5%	<input type="cross"/> Inefficient

4. Response Curves

Non-linear SHAP vs. Spend curves reveal **diminishing returns** — where additional spend yields minimal incremental revenue.

💡 Business Applications

Budget Reallocation

Current Allocation		Recommended Allocation	
TV:	35%	TV:	40% (+5%)
Facebook:	25%	Facebook:	28% (+3%)
Search:	20%	Search:	15% (-5%)
OOH:	12%	OOH:	10% (-2%)
Print:	8%	Print:	7% (-1%)

Scenario Simulation

- **What if** TV spend increases by 20%?
- **What if** we cut Print entirely?
- **What if** we double Search during Q4?

Apply adstock transformation → Pass through XGBoost → Project revenue impact.

⚠ Limitations & Assumptions

Assumption	Implication
Static pricing	Revenue changes attributed to marketing, not price
No macro shocks	External events (COVID, recession) not modeled
Linear adstock decay	May oversimplify complex carryover dynamics
Correlation ≠ Causation	SHAP measures association, not true causal effect
US holidays only	Results may differ for other geographies

🔧 Troubleshooting

Issue	Solution
Prophet installation fails	Use <code>pip install prophet==1.1.4</code> or conda-forge
Memory error during Optuna	Reduce <code>n_trials</code> or <code>n_estimators</code> range
MAPE returns NaN	Check for zero revenue values; use SMAPE instead
SHAP feature mismatch	Ensure train/test have identical columns
Plotly charts not rendering	Set <code>pio.renderers.default = 'notebook'</code>

References

- [Prophet Documentation](#)
 - [XGBoost Documentation](#)
 - [SHAP Documentation](#)
 - [Optuna Documentation](#)
 - [Robyn by Meta \(MMM inspiration\)](#)
-

License

This project is licensed under the **MIT License** — see [LICENSE](#) for details.

MIT License

Copyright (c) 2024 Praveen Kumar Anwla

Permission is hereby granted, free of charge, to any person obtaining a copy of this software and associated documentation files (the "Software"), to deal in the Software without restriction...

Credits

Original Author: Praveen Kumar Anwla

Organization: [TowardsMachineLearning.Org](#)

Contributing

1. Fork the repository
 2. Create a feature branch (`git checkout -b feature/improvement`)
 3. Commit changes (`git commit -m 'Add improvement'`)
 4. Push to branch (`git push origin feature/improvement`)
 5. Open a Pull Request
-

★ Star this repo if you find it useful!

Made with ❤️ for Marketing Analytics