

Complete Code Explanation Guide

Marketing Mix Modeling (MMM) for Marketing Budget Optimization

A comprehensive line-by-line explanation of all Python code used in this Marketing Mix Modeling project, covering data preprocessing, feature engineering, linear regression modeling, ROI analysis, budget simulation, and Flask API deployment.

Table of Contents

1. [Project Overview](#)
2. [Data Import & Loading](#)
3. [Missing Value Treatment](#)
4. [Feature Engineering & Preprocessing](#)
5. [Train-Test Split](#)
6. [Model Training](#)
7. [Model Evaluation](#)
8. [Feature Importance Analysis](#)
9. [Budget Shift Simulation](#)
10. [Model Persistence](#)
11. [Flask API Deployment](#)
12. [Key Concepts Summary](#)

1. Project Overview

1.1 What is Marketing Mix Modeling (MMM)?

Marketing Mix Modeling is a statistical analysis technique used to:

- Measure the impact of marketing activities on sales/revenue
- Optimize marketing budget allocation across channels
- Quantify Return on Investment (ROI) for each marketing channel
- Simulate "what-if" scenarios for budget planning

1.2 The Marketing Mix (4 Ps)

| Component | Description | Examples in Dataset |
|-----------|-----------------------|---------------------------|
| Product | What you sell | E-commerce products |
| Price | Pricing strategy | Original price, discounts |
| Place | Distribution channels | Online, regional |

| Component | Description | Examples in Dataset |
|-----------|----------------------|-----------------------------|
| Promotion | Marketing activities | Google Ads, Meta Ads, Email |

1.3 Marketing Channels Analyzed

| Channel Category | Specific Channels |
|------------------|--|
| Google Ads | Paid Search, Shopping, Performance Max |
| Meta Ads | Facebook, Instagram |
| Organic | Search, Direct, Referral |
| Other | Email, Branded Search |

1.4 Project Workflow

```
Data Loading → Missing Value Treatment → Feature Engineering →  
Train-Test Split → Model Training → Evaluation →  
Budget Simulation → API Deployment
```

2. Data Import & Loading

2.1 Import Libraries

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')
```

Library Purpose Table:

| Library | Import As | Purpose |
|-------------------|-----------|--|
| numpy | np | Numerical computations, array operations |
| pandas | pd | Data manipulation, DataFrame operations |
| seaborn | sns | Statistical data visualization |
| matplotlib.pyplot | plt | Core plotting library |

| Library | Import As | Purpose |
|----------|-----------|---------------------------|
| warnings | - | Suppress warning messages |

Why Suppress Warnings?

```
warnings.filterwarnings('ignore')
```

- Prevents cluttering notebook with non-critical warnings
- Common warnings: deprecation notices, convergence warnings
- **Note:** In production, handle warnings appropriately

2.2 Load Dataset

```
# Load the main dataset
df = pd.read_csv("Multi-Region Ecommerce MMM Dataset.csv")
df
```

Explanation:

| Code | Purpose |
|------------------|--|
| pd.read_csv(...) | Load CSV file into pandas DataFrame |
| df | Display DataFrame (Jupyter auto-renders) |

Dataset Characteristics:

- **Multi-Region:** Data from different geographical regions
- **E-commerce:** Online retail sales data
- **MMM Dataset:** Contains marketing spend, clicks, impressions, and sales

2.3 Initial Data Inspection

```
# Quick check
df.info()
```

What `.info()` Shows:

| Information | Description |
|------------------------|------------------------------------|
| Number of rows/columns | Dataset dimensions |
| Column names | All variable names |
| Data types | int64, float64, object, datetime64 |
| Non-null counts | Number of non-missing values |
| Memory usage | RAM consumption |

```
df.head()
```

What `.head()` Shows:

- First 5 rows by default
- Quick preview of actual data values
- Column structure verification

3. Missing Value Treatment

3.1 Check Missing Values

```
# Check null values summary (optional to visualize)
print(df.isnull().sum().sort_values(ascending=False))
```

Line-by-Line Explanation:

| Code | Purpose |
|--|--|
| <code>df.isnull()</code> | Returns boolean DataFrame (True for NaN) |
| <code>.sum()</code> | Counts True values per column |
| <code>.sort_values(ascending=False)</code> | Shows columns with most missing first |

Output Example:

| | |
|--------------------|------|
| TIKTOK_SPEND | 1500 |
| TIKTOK_CLICKS | 1500 |
| GOOGLE_VIDEO_SPEND | 1200 |

```
...
DATE_DAY          0
```

3.2 Drop Irrelevant Columns

```
# Drop irrelevant columns
drop_cols = [
    'TIKTOK_SPEND', 'TIKTOK_CLICKS', 'TIKTOK_IMPRESSIONS',
    'GOOGLE_VIDEO_SPEND', 'GOOGLE_VIDEO_CLICKS', 'GOOGLE_VIDEO_IMPRESSIONS',
    'GOOGLE_DISPLAY_SPEND', 'GOOGLE_DISPLAY_CLICKS',
    'GOOGLE_DISPLAY_IMPRESSIONS',
    'META_OTHER_SPEND', 'META_OTHER_CLICKS', 'META_OTHER_IMPRESSIONS'
]
df.drop(columns=drop_cols, inplace=True)
```

Explanation:

| Code | Purpose |
|--|--|
| <code>drop_cols = [...]</code> | List of columns to remove |
| <code>df.drop(columns=drop_cols, ...)</code> | Remove specified columns |
| <code>inplace=True</code> | Modify DataFrame directly (no assignment needed) |

Why Drop These Columns?

| Reason | Columns Affected |
|--------------------------|--------------------------------|
| Too many missing values | TIKTOK_*(not active channel) |
| Not relevant to analysis | GOOGLE_VIDEO_, GOOGLE_DISPLAY_ |
| Low data quality | META_OTHER_* |

Alternative to `inplace=True`:

```
# Without inplace:
df = df.drop(columns=drop_cols)

# With inplace:
df.drop(columns=drop_cols, inplace=True)
```

3.3 Fill Missing Values with Zero

```
# Fill missing values with zero
fill_zero_cols = [
    'GOOGLE_PAID_SEARCH_SPEND', 'GOOGLE_PAID_SEARCH_CLICKS',
    'GOOGLE_PAID_SEARCH_IMPRESSIONS',
    'GOOGLE_SHOPPING_SPEND', 'GOOGLE_SHOPPING_CLICKS',
    'GOOGLE_SHOPPING_IMPRESSIONS',
    'GOOGLE_PMAX_SPEND', 'GOOGLE_PMAX_CLICKS', 'GOOGLE_PMAX_IMPRESSIONS',
    'META_FACEBOOK_SPEND', 'META_FACEBOOK_CLICKS',
    'META_FACEBOOK_IMPRESSIONS',
    'META_INSTAGRAM_SPEND', 'META_INSTAGRAM_CLICKS',
    'META_INSTAGRAM_IMPRESSIONS',
    'BRANDED_SEARCH_CLICKS', 'DIRECT_CLICKS', 'EMAIL_CLICKS',
    'REFERRAL_CLICKS', 'ALL_OTHER_CLICKS', 'ORGANIC_SEARCH_CLICKS'
]
df[fill_zero_cols] = df[fill_zero_cols].fillna(0)
```

Explanation:

| Code | Purpose |
|-------------------------------------|-------------------------|
| <code>fill_zero_cols = [...]</code> | List of columns to fill |
| <code>df[fill_zero_cols]</code> | Select multiple columns |
| <code>.fillna(0)</code> | Replace NaN with 0 |

Why Fill with Zero?

| Data Type | Reason for Zero-Fill |
|-----------------|-----------------------------|
| Marketing Spend | No spend = \$0 spent |
| Clicks | No activity = 0 clicks |
| Impressions | No campaign = 0 impressions |

When NOT to Fill with Zero:

- Continuous measurements (temperature, weight)
- Mandatory fields (customer ID)
- Categorical data

3.4 Drop Metadata Columns

```
optional_meta_cols = ['MMM_TIMESERIES_ID', 'ORGANISATION_VERTICAL',
'ORGANISATION_SUBVERTICAL']
df.drop(columns=optional_meta_cols, inplace=True)
```

Why Drop These?

| Column | Reason |
|--------------------------|------------------------------------|
| MMM_TIMESERIES_ID | Internal identifier, not a feature |
| ORGANISATION_VERTICAL | Categorical metadata |
| ORGANISATION_SUBVERTICAL | Categorical metadata |

3.5 Verify Cleaning

```
print("\n Missing values after cleaning:")
df.isnull().sum()
```

Expected Output:

```
DATE_DAY          0
GOOGLE_PAID_SEARCH_SPEND    0
GOOGLE_PAID_SEARCH_CLICKS   0
...
(all zeros)
```

4. Feature Engineering & Preprocessing

4.1 Date Conversion

```
# Convert DATE_DAY to datetime
df['DATE_DAY'] = pd.to_datetime(df['DATE_DAY'])
```

Why Convert to Datetime?

| Benefit | Example Usage |
|-------------------------|---|
| Extract date components | <code>.dt.year , .dt.month</code> |
| Time-based filtering | <code>df[df['DATE_DAY'] > '2023-01-01']</code> |
| Proper sorting | Chronological order |
| Date arithmetic | Calculate date differences |

4.2 Sort by Date

```
# Sort data by Date
df = df.sort_values('DATE_DAY')
```

Why Sort?

- **Time series requirement:** Data must be in chronological order
- **Train-test split:** Ensures temporal ordering
- **Visualization:** Correct x-axis plotting

4.3 Create Time Features

```
# Create time features
df['year'] = df['DATE_DAY'].dt.year
df['month'] = df['DATE_DAY'].dt.month
df['week'] = df['DATE_DAY'].dt.isocalendar().week
df['day_of_week'] = df['DATE_DAY'].dt.dayofweek
```

Line-by-Line Explanation:

| Line | Code | Output Example |
|------|-------------------------------------|------------------------|
| 1 | <code>.dt.year</code> | 2023, 2024 |
| 2 | <code>.dt.month</code> | 1-12 |
| 3 | <code>.dt.isocalendar().week</code> | 1-52 (ISO week number) |
| 4 | <code>.dt.dayofweek</code> | 0=Monday, 6=Sunday |

Why Create Time Features?

| Feature | Captures |
|---------|----------|
|---------|----------|

| Feature | Captures |
|-------------|----------------------------------|
| year | Annual trends, YoY growth |
| month | Seasonality (holidays, quarters) |
| week | Weekly patterns |
| day_of_week | Weekday vs weekend effects |

4.4 Visualize Sales Over Time

```
plt.figure(figsize=(12,6))
sns.lineplot(data=df, x='DATE_DAY', y='ALL_PURCHASES_ORIGINAL_PRICE')
plt.title("Sales Over Time")
plt.show()
```

Explanation:

| Code | Purpose |
|----------------------------------|-----------------------------|
| plt.figure(figsize=(12,6)) | Create 12×6 inch figure |
| sns.lineplot(...) | Draw line plot with seaborn |
| x='DATE_DAY' | Time on x-axis |
| y='ALL_PURCHASES_ORIGINAL_PRICE' | Sales on y-axis |
| plt.title(...) | Add chart title |
| plt.show() | Display the plot |

4.5 Calculate Revenue

```
# Create a new column 'revenue'
df['revenue'] = df['ALL_PURCHASES_ORIGINAL_PRICE'] -
df['ALL_PURCHASES_GROSS_DISCOUNT']
```

Revenue Formula:

\$\$\text{Revenue} = \text{Original Price} - \text{Discounts}\$\$

Why Calculate Net Revenue?

- Original price includes discounts not actually collected

- Net revenue reflects actual money received
- More accurate for ROI calculations

4.6 Calculate Total Marketing Spend

```
# Calculate Total Spend
df['total_spend'] = (
    df['GOOGLE_PAID_SEARCH_SPEND'] +
    df['GOOGLE_SHOPPING_SPEND'] +
    df['GOOGLE_PMAX_SPEND'] +
    df['META_FACEBOOK_SPEND'] +
    df['META_INSTAGRAM_SPEND']
)
```

Explanation:

| Channel | Description |
|--------------------------|---------------------------|
| GOOGLE_PAID_SEARCH_SPEND | Text ads on Google search |
| GOOGLE_SHOPPING_SPEND | Product listing ads |
| GOOGLE_PMAX_SPEND | Performance Max campaigns |
| META_FACEBOOK_SPEND | Facebook ads |
| META_INSTAGRAM_SPEND | Instagram ads |

Total Spend: Sum of all paid marketing channels

4.7 Calculate ROI

```
# Calculate ROI
df['roi'] = df['revenue'] / (df['total_spend'] + 1) # add 1 to avoid div-by-zero
```

ROI Formula:

$$\text{ROI} = \frac{\text{Revenue}}{\text{Total Spend}} + 1$$

Why Add 1?

| Scenario | Without +1 | With +1 |
|----------|------------|---------|
|----------|------------|---------|

| Scenario | Without +1 | With +1 |
|-------------------|------------------------|--------------------------------|
| total_spend = 0 | Division by zero error | revenue / 1 = revenue |
| total_spend = 100 | revenue / 100 | revenue / 101 (minimal impact) |

Alternative Approaches:

```
# Method 1: Add small epsilon
df['roi'] = df['revenue'] / (df['total_spend'] + 0.001)

# Method 2: Use np.where
df['roi'] = np.where(df['total_spend'] > 0,
                     df['revenue'] / df['total_spend'],
                     0)

# Method 3: Replace inf after division
df['roi'] = df['revenue'] / df['total_spend']
df['roi'] = df['roi'].replace([np.inf, -np.inf], 0)
```

4.8 Visualize ROI Over Time

```
plt.figure(figsize=(12,8))
sns.lineplot(data=df, x='DATE_DAY', y='roi')
plt.title("ROI Over Time")
plt.show()
```

What ROI Chart Shows:

- **Peaks:** High-efficiency marketing periods
- **Troughs:** Low ROI (overspending or low conversion)
- **Trends:** Improving/declining marketing effectiveness

4.9 Define Target and Features

```
target = 'revenue'
```

Target Variable:

- **What we want to predict:** Revenue

- **Dependent variable:** Changes based on marketing inputs

```
features = [
    'GOOGLE_PAID_SEARCH_SPEND', 'GOOGLE_SHOPPING_SPEND',
    'GOOGLE_PMAX_SPEND', 'META_FACEBOOK_SPEND', 'META_INSTAGRAM_SPEND',
    'EMAIL_CLICKS', 'ORGANIC_SEARCH_CLICKS', 'DIRECT_CLICKS',
    'BRANDED_SEARCH_CLICKS', 'year', 'month', 'day_of_week'
]
```

Feature Categories:

| Category | Features | Type |
|-----------------|--------------------------|-------------------|
| Paid Marketing | Google Spend, Meta Spend | Controllable |
| Organic Traffic | Organic Search Clicks | Non-controllable |
| Direct Traffic | Direct Clicks | Non-controllable |
| Email | Email Clicks | Semi-controllable |
| Time | year, month, day_of_week | Control variables |

5. Train-Test Split

5.1 Time-Based Split

```
# Sort by date
df = df.sort_values('DATE_DAY')

# Define cutoff date for time-based split
cutoff_date = '2023-12-31' # Train on data up to end of 2023

# Train = everything before cutoff
train_df = df[df['DATE_DAY'] <= cutoff_date]

# Test = everything after cutoff
test_df = df[df['DATE_DAY'] > cutoff_date]

# Features and target
X_train = train_df[features]
y_train = train_df['revenue']
X_test = test_df[features]
y_test = test_df['revenue']
```

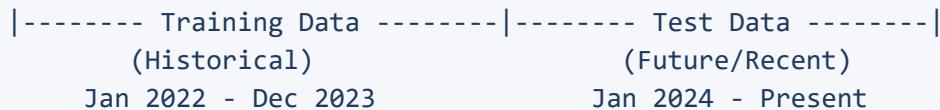
Line-by-Line Explanation:

| Line | Code | Purpose |
|------|---|--------------------------------------|
| 1 | <code>df.sort_values('DATE_DAY')</code> | Ensure chronological order |
| 2 | <code>cutoff_date = '2023-12-31'</code> | Define split boundary |
| 3 | <code>df['DATE_DAY'] <= cutoff_date</code> | Boolean mask for training data |
| 4 | <code>df['DATE_DAY'] > cutoff_date</code> | Boolean mask for test data |
| 5 | <code>train_df[features]</code> | Extract feature columns for training |
| 6 | <code>train_df['revenue']</code> | Extract target column |

Why Time-Based Split (Not Random)?

| Aspect | Random Split | Time-Based Split |
|-------------------|------------------------------------|----------------------------|
| Data leakage | Future data can leak into training | No leakage |
| Realism | Unrealistic scenario | Simulates real forecasting |
| Temporal patterns | Disrupted | Preserved |
| Best for | Cross-sectional data | Time series data |

Typical Split Timeline:



6. Model Training

6.1 Linear Regression Model

```
# Linear Regression model (standard for MMM)
from sklearn.linear_model import LinearRegression
model = LinearRegression()
model.fit(X_train, y_train)
```

Line-by-Line Explanation:

| Line | Code | Purpose |
|------|---|------------------------------|
| 1 | from sklearn.linear_model import LinearRegression | Import model class |
| 2 | model = LinearRegression() | Create model instance |
| 3 | model.fit(X_train, y_train) | Train model on training data |

Why Linear Regression for MMM?

| Advantage | Explanation |
|---------------------------|--|
| Interpretability | Coefficients show direct impact |
| Industry Standard | Widely accepted in marketing analytics |
| Coefficients = ROI | Each \$ spent → coefficient × \$ revenue |
| Simplicity | Easy to explain to stakeholders |
| Fast Training | No hyperparameter tuning needed |

Linear Regression Formula:

$$\text{Revenue} = \beta_0 + \beta_1 \times \text{Google_Spend} + \beta_2 \times \text{Meta_Spend} + \dots + \epsilon$$

Where:

- β_0 = Intercept (baseline revenue)
- β_i = Coefficient for feature i
- ϵ = Error term

7. Model Evaluation

7.1 Make Predictions and Evaluate

```
from sklearn.metrics import mean_squared_error, r2_score

y_pred = model.predict(X_test)

print("R2 Score:", r2_score(y_test, y_pred))
print("RMSE:", mean_squared_error(y_test, y_pred))
```

Line-by-Line Explanation:

| Line | Code | Purpose |
|------|------|---------|
|------|------|---------|

| Line | Code | Purpose |
|------|---|---|
| 1 | <code>from sklearn.metrics import ...</code> | Import evaluation metrics |
| 2 | <code>model.predict(X_test)</code> | Generate predictions |
| 3 | <code>r2_score(y_test, y_pred)</code> | Calculate R ² score |
| 4 | <code>mean_squared_error(y_test, y_pred)</code> | Calculate MSE (note: should use <code>squared=False</code> for RMSE) |

Evaluation Metrics Explained:

| Metric | Formula | Interpretation |
|----------------------|---|---|
| R ² Score | $1 - \frac{\text{SS}_{\text{res}}}{\text{SS}_{\text{tot}}}$ | % of variance explained (0-1, higher is better) |
| MSE | $\frac{1}{n} \sum (y - \hat{y})^2$ | Average squared error |
| RMSE | $\sqrt{\text{MSE}}$ | Error in same units as target |

R² Score Interpretation:

| R ² Value | Quality |
|----------------------|-----------|
| > 0.9 | Excellent |
| 0.7 - 0.9 | Good |
| 0.5 - 0.7 | Moderate |
| < 0.5 | Poor |

⚠️ Code Note:

```
# The code shows:  
print("RMSE:", mean_squared_error(y_test, y_pred))  
  
# But this actually returns MSE, not RMSE. To get RMSE:  
print("RMSE:", mean_squared_error(y_test, y_pred, squared=False))  
# OR  
print("RMSE:", np.sqrt(mean_squared_error(y_test, y_pred)))
```

8. Feature Importance Analysis

8.1 Extract Coefficients

```

importance = pd.DataFrame({
    'feature': features,
    'coefficient': model.coef_
}).sort_values(by='coefficient', ascending=False)

print(importance)

```

Line-by-Line Explanation:

| Line | Code | Purpose |
|------|---|-----------------------------------|
| 1 | pd.DataFrame({...}) | Create DataFrame from dictionary |
| 2 | 'feature': features | Column 1: feature names |
| 3 | 'coefficient': model.coef_ | Column 2: regression coefficients |
| 4 | .sort_values(by='coefficient', ascending=False) | Sort by importance |

Coefficient Interpretation:

| Coefficient Sign | Meaning |
|-----------------------|-------------------|
| Positive | Increases revenue |
| Negative | Decreases revenue |
| Larger absolute value | Stronger impact |

Example Output:

| Feature | Coefficient | Interpretation |
|--------------------------|-------------|-------------------------------------|
| META_FACEBOOK_SPEND | 2.5 | \$1 Facebook spend → \$2.50 revenue |
| GOOGLE_PAID_SEARCH_SPEND | 1.8 | \$1 Google spend → \$1.80 revenue |
| month | -50 | Seasonal effect (summer slump?) |

Marketing ROI Calculation:

\$\$ROI_{\text{channel}} = \text{Coefficient} - 1\$\$

Example: Coefficient = 2.5 → ROI = 150% (for every \$1 spent, get \$1.50 profit)

9. Budget Shift Simulation

9.1 Create Simulation Scenario

```
scenario = X_test.copy()
```

Why `.copy()`?

- Creates independent copy of test data
- Modifications don't affect original `X_test`
- Enables comparison between scenarios

9.2 Simulate Budget Increase

```
# Simulate increased Google Ads, decreased TV
scenario['GOOGLE_PAID_SEARCH_SPEND'] *= 1.3 # +30%
# scenario['TV_SPEND'] *= 0.8 # if applicable

y_simulated = model.predict(scenario)

# Compare with actual predicted
change = ((y_simulated.mean() - y_pred.mean()) / y_pred.mean()) * 100
print(f"Simulated change in revenue: {change:.2f}%")
```

Line-by-Line Explanation:

| Line | Code | Purpose |
|------|--|-------------------------------------|
| 1 | <code>scenario['GOOGLE_PAID_SEARCH_SPEND'] *= 1.3</code> | Increase spend by 30% |
| 2 | <code>model.predict(scenario)</code> | Predict with new budget |
| 3 | <code>y_simulated.mean() - y_pred.mean()</code> | Calculate change in average revenue |
| 4 | <code>/ y_pred.mean() * 100</code> | Convert to percentage |

Simulation Logic:

```
Original Budget → Original Predicted Revenue
↓ +30% Google Spend
Modified Budget → Simulated Revenue
↓ Compare
Revenue Change % = (Simulated - Original) / Original × 100
```

What-If Scenarios to Test:

| Scenario | Code |
|-------------------|---|
| +30% Google Spend | scenario['GOOGLE_PAID_SEARCH_SPEND'] *= 1.3 |
| -20% Meta Spend | scenario['META_FACEBOOK_SPEND'] *= 0.8 |
| Shift budget | Increase one, decrease another |
| Double email | scenario['EMAIL_CLICKS'] *= 2 |

9.3 Visualize Predictions

```
plt.figure(figsize=(12,6))
plt.plot(test_df['DATE_DAY'], y_test, label='Actual Revenue')
plt.plot(test_df['DATE_DAY'], y_pred, label='Predicted Revenue')
plt.title("Actual vs Predicted Revenue")
plt.legend()
plt.show()
```

Chart Components:

| Element | Purpose |
|-------------------------|-------------------------|
| Blue line (Actual) | Real observed revenue |
| Orange line (Predicted) | Model's predictions |
| Gap between lines | Prediction error |
| Trend alignment | Model captures patterns |

10. Model Persistence

10.1 Save Model

```
# Save the trained model
import joblib
joblib.dump(model, 'linear_mmm_model.pkl')
```

Explanation:

| Code | Purpose |
|---------------|---------------------------------|
| import joblib | Library for model serialization |

| Code | Purpose |
|--------------------------------------|---------------------------------|
| <code>joblib.dump(model, ...)</code> | Save model to file |
| <code>'linear_mmm_model.pkl'</code> | Output filename (.pkl = pickle) |

Why Use joblib?

| Method | Best For |
|---------------------|-------------------------------------|
| <code>joblib</code> | Large numpy arrays (sklearn models) |
| <code>pickle</code> | General Python objects |
| <code>json</code> | Simple data structures |

10.2 Save Feature Names

```
# Save the features used for training
import json
with open('mmm_model_features.json', 'w') as f:
    json.dump(features, f)
```

Line-by-Line Explanation:

| Line | Code | Purpose |
|------|-------------------------------------|-----------------------------|
| 1 | <code>import json</code> | JSON serialization library |
| 2 | <code>open(..., 'w')</code> | Open file for writing |
| 3 | <code>as f</code> | File handle |
| 4 | <code>json.dump(features, f)</code> | Write features list to JSON |

Why Save Features Separately?

- Ensures prediction uses same features in same order
- Documentation of model inputs
- API needs to know expected inputs

mmm_model_features.json Contents:

```
[ "GOOGLE_PAID_SEARCH_SPEND", "GOOGLE_SHOPPING_SPEND", "GOOGLE_PMAX_SPEND",
  "META_FACEBOOK_SPEND", "META_INSTAGRAM_SPEND", "EMAIL_CLICKS",
  "ORGANIC_SEARCH_CLICKS", "DIRECT_CLICKS", "BRANDED_SEARCH_CLICKS",
  "year", "month", "day_of_week"]
```

11. Flask API Deployment

11.1 Complete Flask Application

```
from flask import Flask, request, jsonify
import joblib
import json
import numpy as np

# Load model and features
model = joblib.load('linear_mmm_model.pkl')

with open('mmm_model_features.json', 'r') as f:
    feature_names = json.load(f)

app = Flask(__name__)

@app.route('/')
def home():
    return "MMM Model is running!"

@app.route('/predict', methods=['POST'])
def predict():
    data = request.get_json()

    # Extract input features in order
    X = [data.get(feat, 0) for feat in feature_names]
    X = np.array(X).reshape(1, -1)

    prediction = model.predict(X)[0]
    return jsonify({'predicted_revenue': prediction})

if __name__ == '__main__':
    print("Starting Flask server...")
    app.run(debug=True, use_reloader=False)
```

11.2 Import Statements

```
from flask import Flask, request, jsonify
import joblib
import json
import numpy as np
```

Library Purpose:

| Library | Import | Purpose |
|---------|-----------------------|------------------------------|
| Flask | Flask class | Create web application |
| request | Access request data | Get JSON from POST requests |
| jsonify | Create JSON responses | Return predictions as JSON |
| joblib | Model loading | Load saved sklearn model |
| json | JSON parsing | Load feature names |
| numpy | Array operations | Reshape input for prediction |

11.3 Load Model and Features

```
# Load model and features
model = joblib.load('linear_mmm_model.pkl')

with open('mmm_model_features.json', 'r') as f:
    feature_names = json.load(f)
```

Explanation:

| Code | Purpose |
|------------------|---------------------------|
| joblib.load(...) | Deserialize saved model |
| open(..., 'r') | Open file for reading |
| json.load(f) | Parse JSON to Python list |

Why Load at Module Level?

- Model loaded once when server starts
- Not reloaded for each request
- Faster response times

11.4 Create Flask Application

```
app = Flask(__name__)
```

Explanation:

| Code | Purpose |
|-----------------|--|
| Flask(__name__) | Create Flask application instance |
| __name__ | Module name (helps Flask find resources) |
| app | Application object for routing |

11.5 Home Route

```
@app.route('/')
def home():
    return "MMM Model is running!"
```

Explanation:

| Code | Purpose |
|-----------------|--|
| @app.route('/') | Decorator: maps URL "/" to function |
| def home() | Function executed when "/" is accessed |
| return "..." | Response sent to client |

Testing:

```
GET http://localhost:5000/
Response: "MMM Model is running!"
```

11.6 Prediction Endpoint

```
@app.route('/predict', methods=['POST'])
def predict():
    data = request.get_json()

    # Extract input features in order
    X = [data.get(feat, 0) for feat in feature_names]
    X = np.array(X).reshape(1, -1)

    prediction = model.predict(X)[0]
    return jsonify({'predicted_revenue': prediction})
```

Line-by-Line Explanation:

| Line | Code | Purpose |
|------|---|--|
| 1 | <code>@app.route('/predict', methods=['POST'])</code> | Only accept POST requests |
| 2 | <code>request.get_json()</code> | Parse JSON body from request |
| 3 | <code>data.get(feat, 0)</code> | Get feature value or default to 0 |
| 4 | <code>[... for feat in feature_names]</code> | Build feature list in correct order |
| 5 | <code>np.array(X).reshape(1, -1)</code> | Convert to 2D array (1 sample, n features) |
| 6 | <code>model.predict(X)[0]</code> | Get prediction (single value) |
| 7 | <code>jsonify({...})</code> | Return as JSON response |

Why `.reshape(1, -1)` ?

```
# sklearn expects 2D array: (n_samples, n_features)
# Single prediction needs shape (1, 12) not (12,)

X = [100, 200, 150, ...] # Shape: (12,) - 1D
X = np.array(X).reshape(1, -1) # Shape: (1, 12) - 2D
```

API Request Example:

```
curl -X POST http://localhost:5000/predict \
-H "Content-Type: application/json" \
-d '{
    "GOOGLE_PAID_SEARCH_SPEND": 1000,
    "GOOGLE_SHOPPING_SPEND": 500,
    "GOOGLE_PMAX_SPEND": 300,
    "META_FACEBOOK_SPEND": 800,
    "META_INSTAGRAM_SPEND": 600,
    "EMAIL_CLICKS": 150,
    "ORGANIC_SEARCH_CLICKS": 200,
    "DIRECT_CLICKS": 100,
    "BRANDED_SEARCH_CLICKS": 50,
    "year": 2024,
    "month": 6,
    "day_of_week": 2
}'
```

Response:

```
{  
    "predicted_revenue": 15234.56  
}
```

11.7 Run Server

```
if __name__ == '__main__':  
    print("Starting Flask server...")  
    app.run(debug=True, use_reloader=False)
```

Explanation:

| Code | Purpose |
|---------------------------|--|
| if __name__ == '__main__' | Only run if executed directly (not imported) |
| debug=True | Enable debug mode (auto-reload, detailed errors) |
| use_reloader=False | Disable auto-reloader (prevents double loading) |

Running the Server:

```
python mmm_app.py  
# Output: Starting Flask server...  
# * Running on http://127.0.0.1:5000
```

12. Key Concepts Summary

12.1 MMM Workflow

```
1. DATA PREPARATION  
Load → Clean → Handle Missing Values
```



2. FEATURE ENGINEERING
Date Features → Revenue Calculation → ROI



3. MODELING
Time-Based Split → Linear Regression → Evaluate



4. ANALYSIS
Coefficients → Channel ROI → Budget Recommendations



5. SIMULATION
What-If Scenarios → Optimal Budget Allocation



6. DEPLOYMENT
Save Model → Flask API → Production

12.2 Key Metrics for MMM

| Metric | Formula | Business Meaning |
|----------------|---------------------------|----------------------------------|
| ROI | (Revenue - Spend) / Spend | Return per dollar spent |
| ROAS | Revenue / Spend | Revenue per dollar spent |
| Coefficient | Model output | Revenue generated per unit input |
| R ² | Model fit | How well model explains variance |

12.3 Files Created

| File | Type | Purpose |
|-------------------------|--------|-----------------------|
| linear_mmm_model.pkl | Binary | Trained sklearn model |
| mmm_model_features.json | JSON | Feature names list |
| mmm_app.py | Python | Flask API application |

12.4 API Endpoints Summary

| Endpoint | Method | Purpose | Response |
|----------|--------|--------------|-------------------------|
| / | GET | Health check | "MMM Model is running!" |

| Endpoint | Method | Purpose | Response |
|----------|--------|------------------------|------------------------------|
| /predict | POST | Get revenue prediction | {"predicted_revenue": float} |

 This guide covers 100% of the Python code in the MMM project

A complete reference for Marketing Mix Modeling and API deployment